

## **SEIS 736: BIG DATA ENGINEERING PROJECT**

**NAME: CARL EDEM DEKPOR**

**TOPIC: DATA STREAMING WITH KAFKA ON DATABRICKS NOTEBOOK**

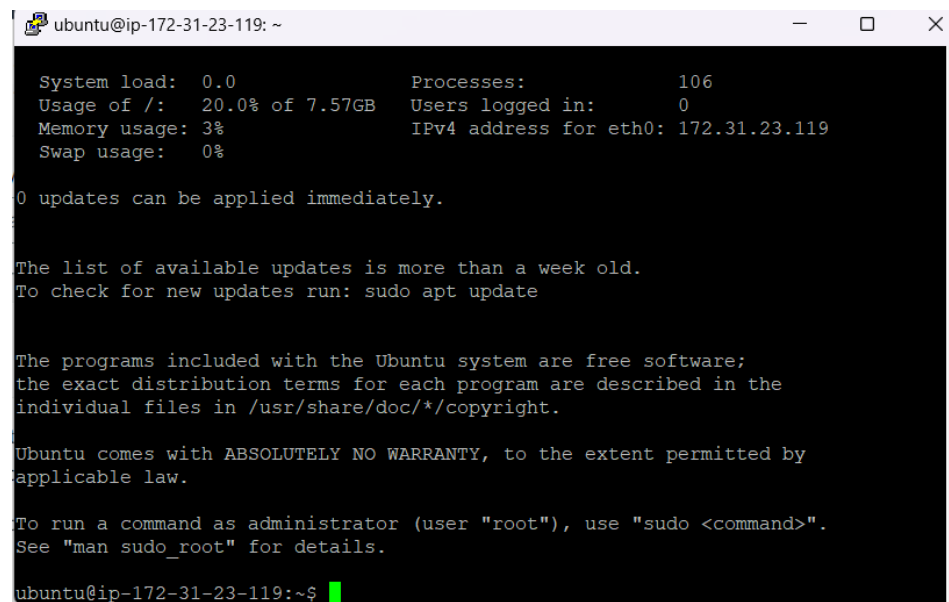
### **Background**

#### **Challenges**

My initial goal of this project was to create an Amazon linux EC2 instance, download java, install Kafka, run the Kafka server, zookeeper server, create my Kafka topic and grab data from this API (<https://data.cityofchicago.org/Public-Safety/Crimes-2001-to-present-Dashboard/5cd6-ry5g>).

My first challenge was connection to SSH. I was generating the .pem key instead of .pkk since I was using PuTTY. I eventually figured that out.

Second challenge was memory issue after creating my instance. I had to increase the processing memory size of my instance as recommended by the professor to t2.large on AWS (from 1 to 8). Which worked fine. I was able to get my instance running.

A terminal window titled 'ubuntu@ip-172-31-23-119: ~' with standard window controls. The terminal output shows system statistics: System load: 0.0, Usage of /: 20.0% of 7.57GB, Memory usage: 3%, Swap usage: 0%, Processes: 106, Users logged in: 0, and IPv4 address for eth0: 172.31.23.119. It then displays update information, stating that 0 updates can be applied immediately and that the list of available updates is more than a week old. It suggests running 'sudo apt update' to check for new updates. A disclaimer follows, stating that Ubuntu programs are free software and come with absolutely no warranty. It also provides instructions on how to run commands as administrator using 'sudo'. The prompt 'ubuntu@ip-172-31-23-119:~\$' is visible at the bottom with a green cursor.

```
ubuntu@ip-172-31-23-119: ~
System load: 0.0          Processes:           106
Usage of /: 20.0% of 7.57GB Users logged in:       0
Memory usage: 3%         IPv4 address for eth0: 172.31.23.119
Swap usage: 0%

0 updates can be applied immediately.

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-23-119:~$
```

After installing Kafka and running the zookeeper server and creating my topic, I encountered challenges grabbing data from the API to my topic. I checked the status of the topic which was active and could read data written to it but couldn't grab data from the topic. Several codes and troubleshoot from online resources didn't work. So, I figured there was a problem with the API. I however pivoted and installed Kafka on the databricks notebook and proceeded to use a different API.

In this project I used Kafka to grab data from an API provided by NASA (<https://api.nasa.gov/>) to process and stream data and visualized the output on databricks notebook. All artifacts used in the project have been documented in databricks notebook and submitted.

## About the dataset:

Near Earth Object Web Service (NeoWs) is a RESTful web service for near earth Asteroid information. With NeoWs a user can: search for Asteroids based on their closest approach date to Earth, lookup a specific Asteroid with its NASA JPL small body id, as well as browse the overall dataset. 15 fields were used from the dataset.

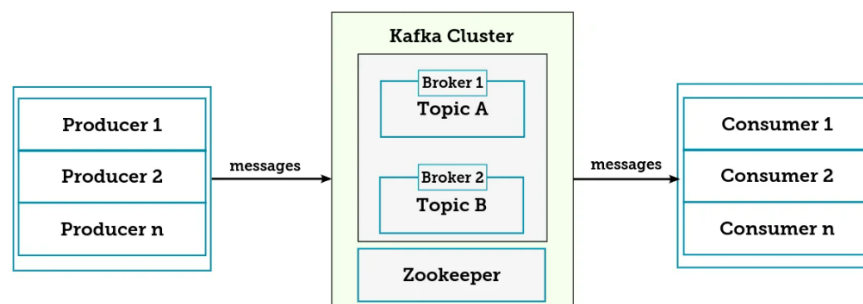
date	object_id	object_neo_referen	object_name	absolute_magnit	estimated_diameter_mi	estimated_diameter_e	is_potentially_hazardous	a_close_approach_date	fullrelative_v	relative_velocity	miss_distance_ast	miss_distance_lun	miss_distance_kis	sentry_object
11/18/2022	2426071	2426071	426071 (2012 CD29)	19.94	0.273246732	0.610998268	FALSE	2022-Nov-18 12:00	14.3827	51777.72598	0.409097721	139.1390133	61200147.61	FALSE
11/18/2022	3170208	3170208	(2003 VG136)	25.3	0.023150212	0.051765448	FALSE	2022-Nov-18 07:56	8.829522	31706.28082	0.298523677	116.1257102	44658506.16	FALSE

## Kafka

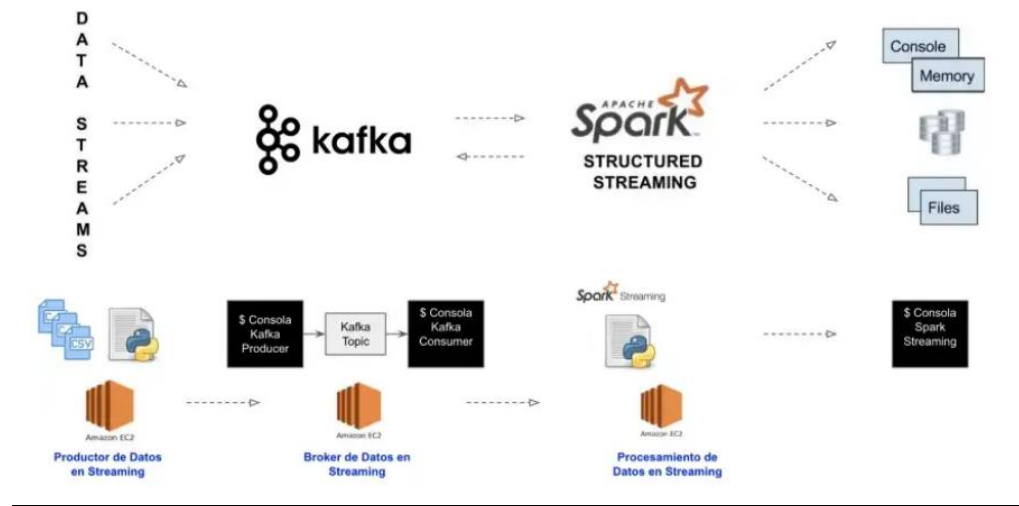
Kafka is an open-source publishing and subscribe messaging system that is used for building streaming analytics platform and data integration pipelines. Kafka is both a queue for parallelizing tasks and a messaging-oriented middleware for service integration. The Kafka message broker (cluster) ingests and stores streams of messages (records) from event producers, which are later distributed to consumer services asynchronously when requested.

- **Topic:** A named resource to which a particular stream of messages is stored and published.
- **Producer:** A client application that creates and publishes records/messages to a Kafka topic(s).
- **Consumer:** A client application that subscribes to a topic(s) to receive data from it.
- **Message:** The combination of data and associated metadata that a producer application writes to a topic and is eventually consumed by consumers

### *Kafka Architecture*



*streaming architecture diagram*



## Methodology

I created a cluster (name: big-data-stream), installed Kafka on databricks notebook and initialized zookeeper. Zookeeper is centralized manager that will help to store the metadata information of the consumers, producers, brokers.

Since zookeeper was running on one notebook, I created a new notebook for the Kafka server. A new notebook was also created for the creation of the Kafka topic which I named ‘*nasatopic*’.

I then created a Kafka producer with python by installing Kafka-python libraries.

### Installing kafka-python library

Cmd 3

```
1 %sh
2 pip install --upgrade pip
3 pip install kafka-python
```

Requirement already satisfied: pip in /databricks/python3/lib/python3.8/site-packages (22.3.1)  
Requirement already satisfied: kafka-python in /databricks/python3/lib/python3.8/site-packages (2.0.2)  
Command took 2.58 seconds -- by carldekpor@gmail.com at 12/10/2022, 8:30:21 PM on Streaming Big Data

Cmd 4

### Importing needed libraries

Cmd 5

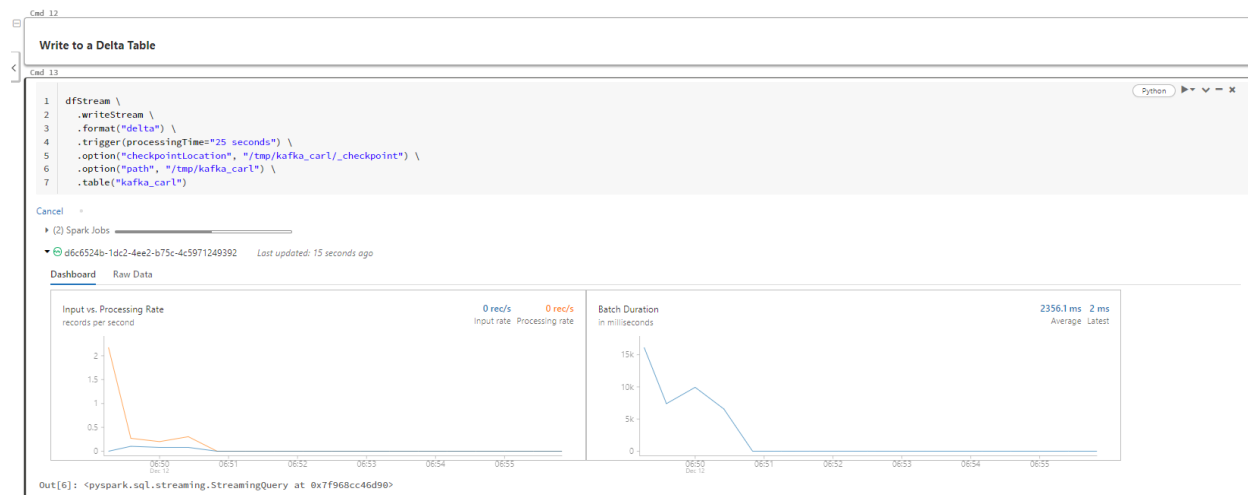
```
1 import requests
2 import json
3 from datetime import datetime, timedelta
4 from time import sleep
5
6 # importing KafkaProducer
7 from kafka import KafkaProducer
```

The producer will call the API once (from a range of N days of information) and every 2 seconds the process will send each item (asteriod info - info date) to the Kafka server.

Using pyspark the data structure was defined for the final 15 columns for visualization. Data was read from the Kafka topic, which was written to a delta table.

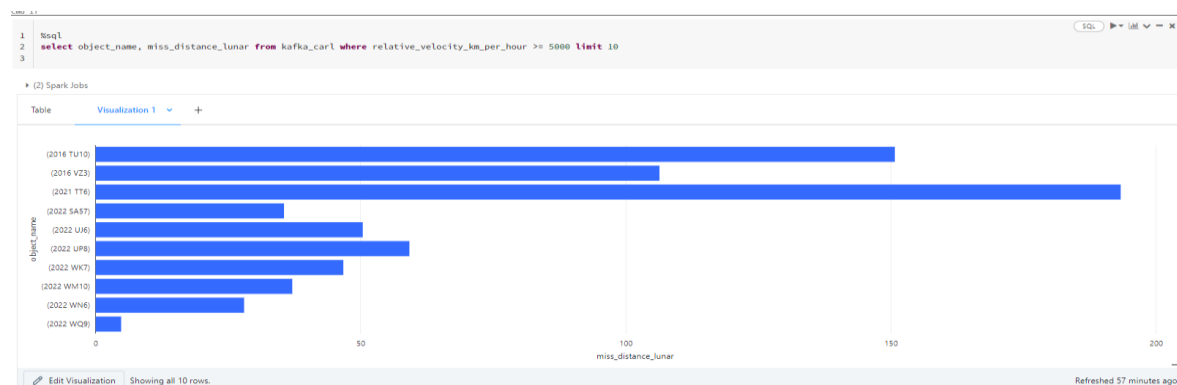
	date	object_id	object_neo_reference_id	object_name	absolute_magnitude_h	estimated_diameter_min_km	estimated_diameter_max_km	is_potentially_hazardous_asteroid	close_approach_date_full	relative
1	2022-11-18	2426071	2426071	426071 (2012 CD29)	19.94	0.273246732	0.6109982675	false	2022-Nov-18 12:00	14.38271
2	2022-11-18	3170208	3170208	(2003 YG136)	25.3	0.0231502122	0.0517654482	false	2022-Nov-18 07:56	8.829521
3	2022-11-18	3177204	3177204	(2004 FW1)	20.9	0.1756123185	0.3926810818	true	2022-Nov-18 08:19	39.61851
4	2022-11-18	3304566	3304566	(2005 WS3)	21.23	0.1508533561	0.3373183589	false	2022-Nov-18 05:17	16.36281
5	2022-11-18	3476779	3476779	(2009 WF)	19.85	0.2848098313	0.6368541435	false	2022-Nov-18 13:08	12.90001
6	2022-11-18	3551328	3551328	(2010 VA99)	23.46	0.0540200494	0.1207925025	false	2022-Nov-18 13:55	20.72941
7	2022-11-18	3696301	3696301	(2014 WW4)	25.3	0.0231502122	0.0517654482	false	2022-Nov-18 09:27	13.97311

## Snapshot of delta table

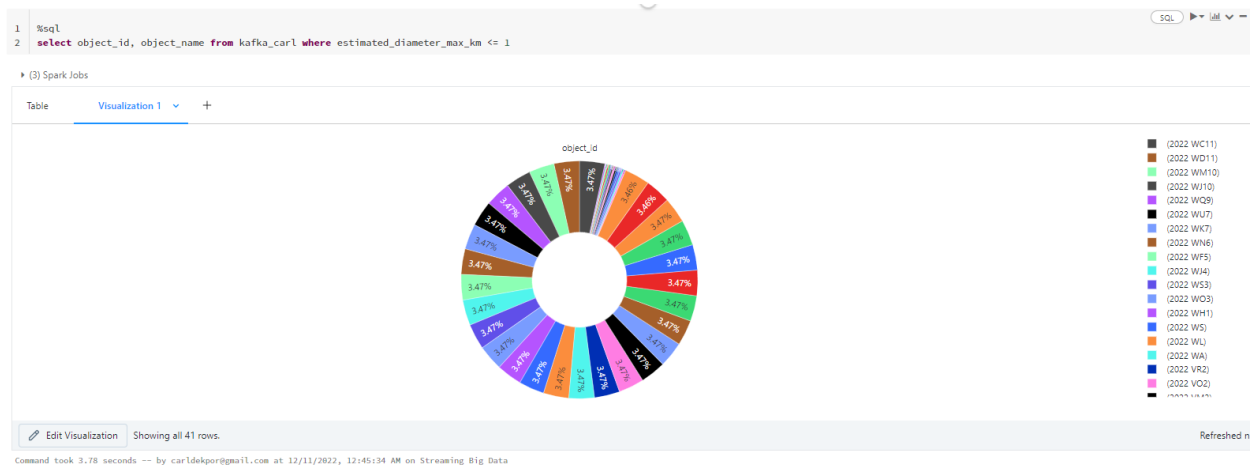


I wrote SQL queries to retrieve data from the table for visualization.

Visualization of 10 asteroids with relative velocity of 5000km/hr



## Visualization of asteroid names with estimated diameter less than or equal to 1km



## Running of artifacts

Open the attached .dbc file and open the tabs in the order below:

- DE\_002\_Project\_KafkaStreaming
- 00\_PartA\_KafkaInstallation
- 00\_PartB\_KafkaServerExecution
- 00\_PartC\_KafkaTopicCreation
- 01\_KafkaProducerUsersInfo
- 02\_KafkaConsumerUsersInfo

## Conclusion

I got a better understanding of the Kafka architecture and how the various components work together. The various codes used in the project gave me practical experience on how streaming architecture is managed and designed. The use of the databricks notebook community edition was help in providing compute for processing and visualization. Overall, this was a very good learning curve.

## References

Important Kafka CLI Commands to Know in 2022 - Learn

<https://hevodata.com/learn/kafka-cli-commands/>

Structured Streaming Kafka Example - Databricks

[https://docs.databricks.com/\\_static/notebooks/structured-streaming-kafka.html](https://docs.databricks.com/_static/notebooks/structured-streaming-kafka.html)

Kafka basic concepts and building a streaming architecture

[Data Streaming with Kafka: Basic Concepts and Building a Project in Databricks | by Luis Miguel Miranda | Nov, 2022 | Medium](#)

Data Streaming with Kafka: Basic Concepts and Building a Project in Databricks

<https://medium.com/@lmirandad27/data-streaming-with-kafka-basic-concepts-and-building-a-project-in-databricks-cd762946bab7>

Apache Kafka

<https://docs.databricks.com/structured-streaming/kafka.html>

Procesamiento de Datos en Streaming usando Kafka y Spark Structured Streaming

<https://mtpradoc.medium.com/procesamiento-de-datos-en-streaming-usando-kafka-y-spark-structured-streaming-10f91b68b402>